



中山大學
SUN YAT-SEN UNIVERSITY



多核程序设计与实践

OpenMP入门

陶钧

taoj23@mail.sysu.edu.cn

中山大学 数据科学与计算机学院
国家超级计算广州中心

- 什么是并行?
- 什么是OpenMP?
- 语法
- 同步机制
- 变量作用域
- 线程调度



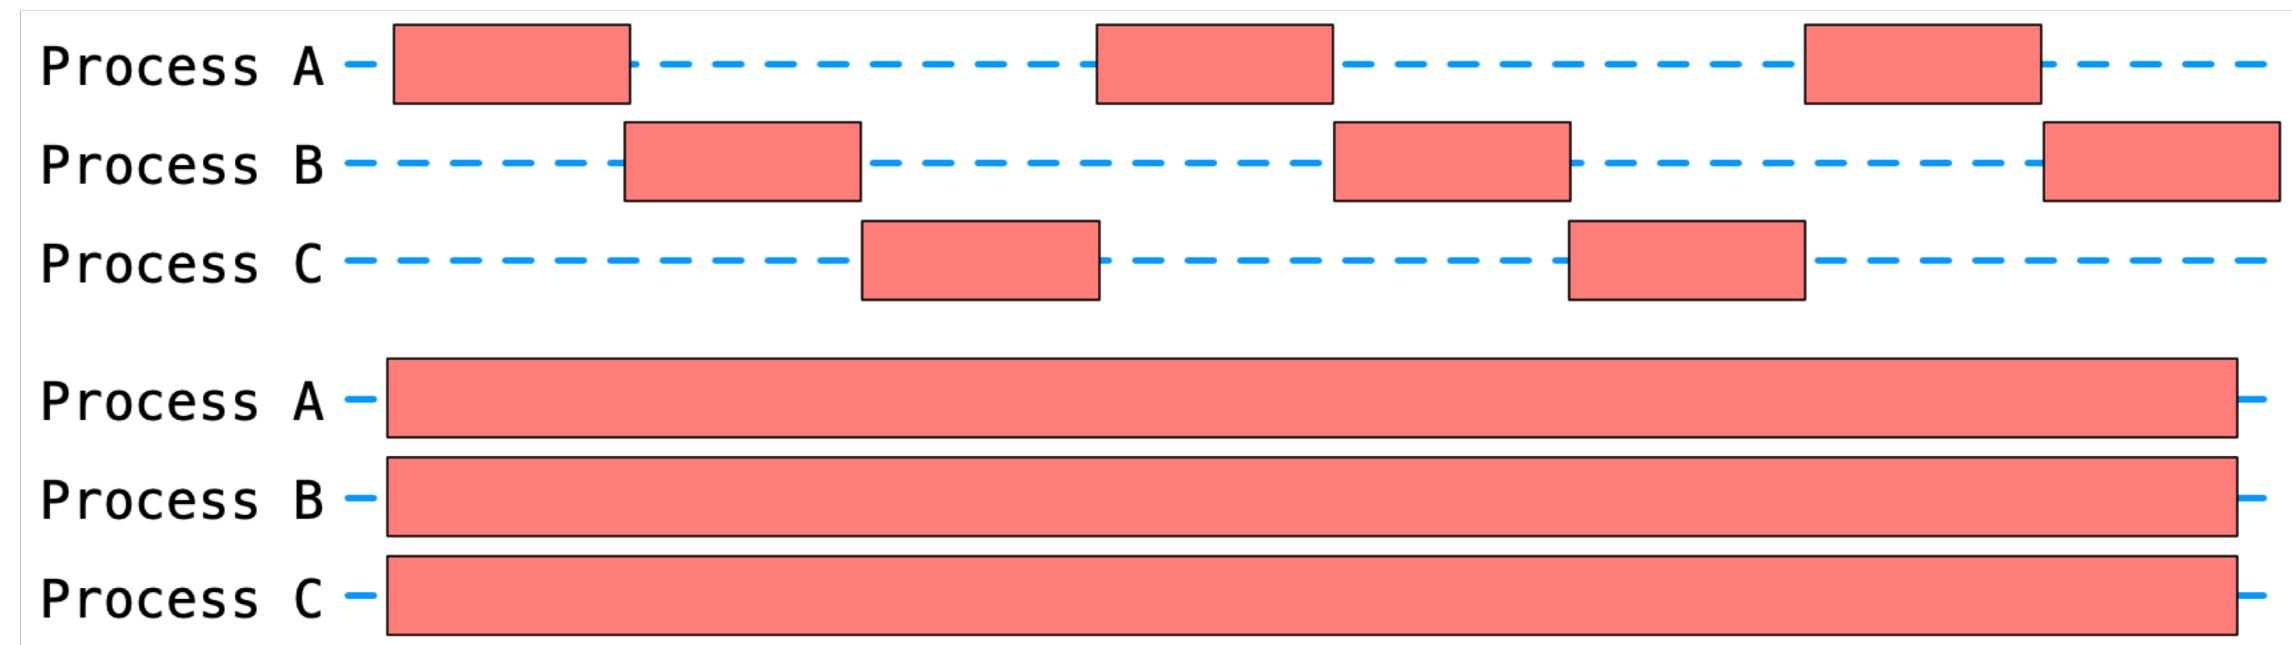
- 多个任务同时进行

- 生活中的多任务与并行

- 边吃饭边看电视边聊天，边听音乐边回邮件
 - 边上课边睡觉，边开车边发短信
 - 写论文、写代码、完成作业（？）

- 计算机上的多任务与并行

交错并发
(interleaved)

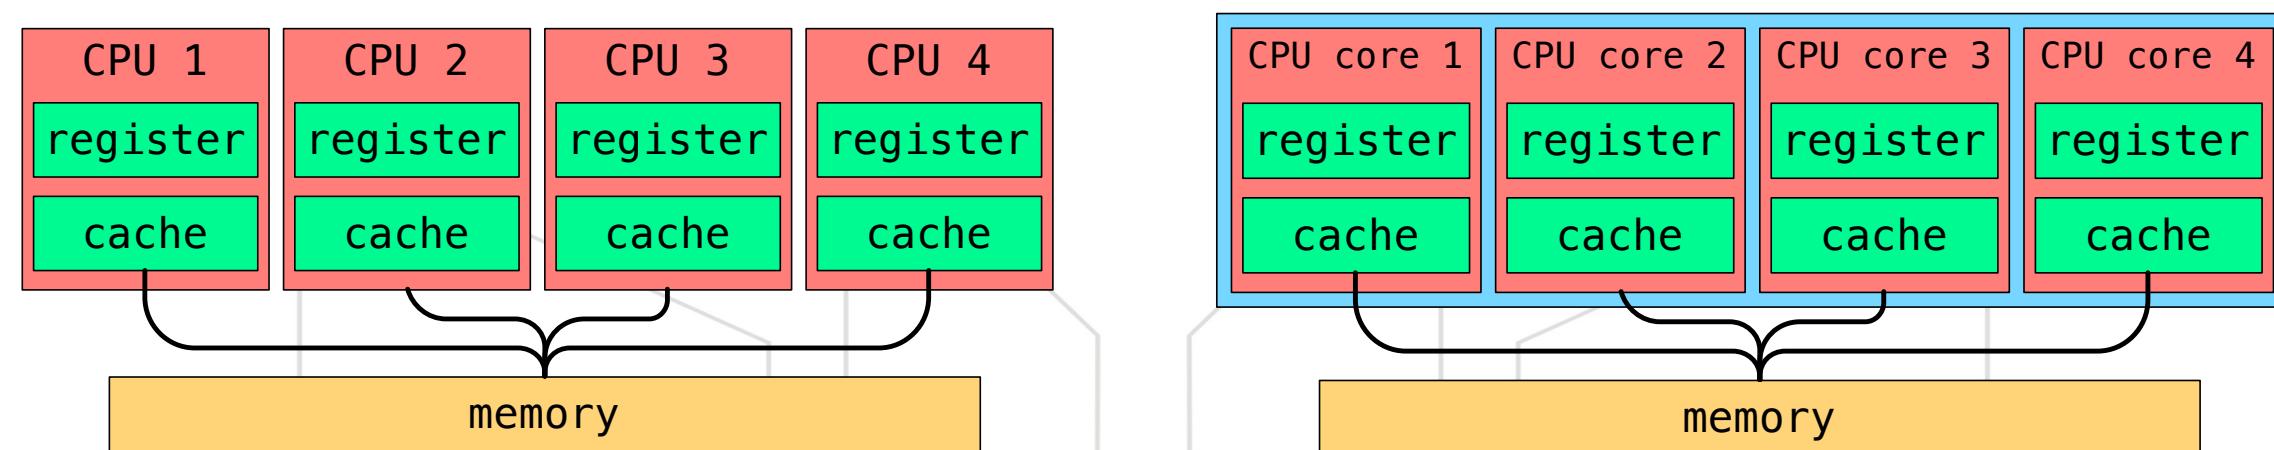


并行
(parallel)

● 多任务实现方式

- 并发 (concurrent)

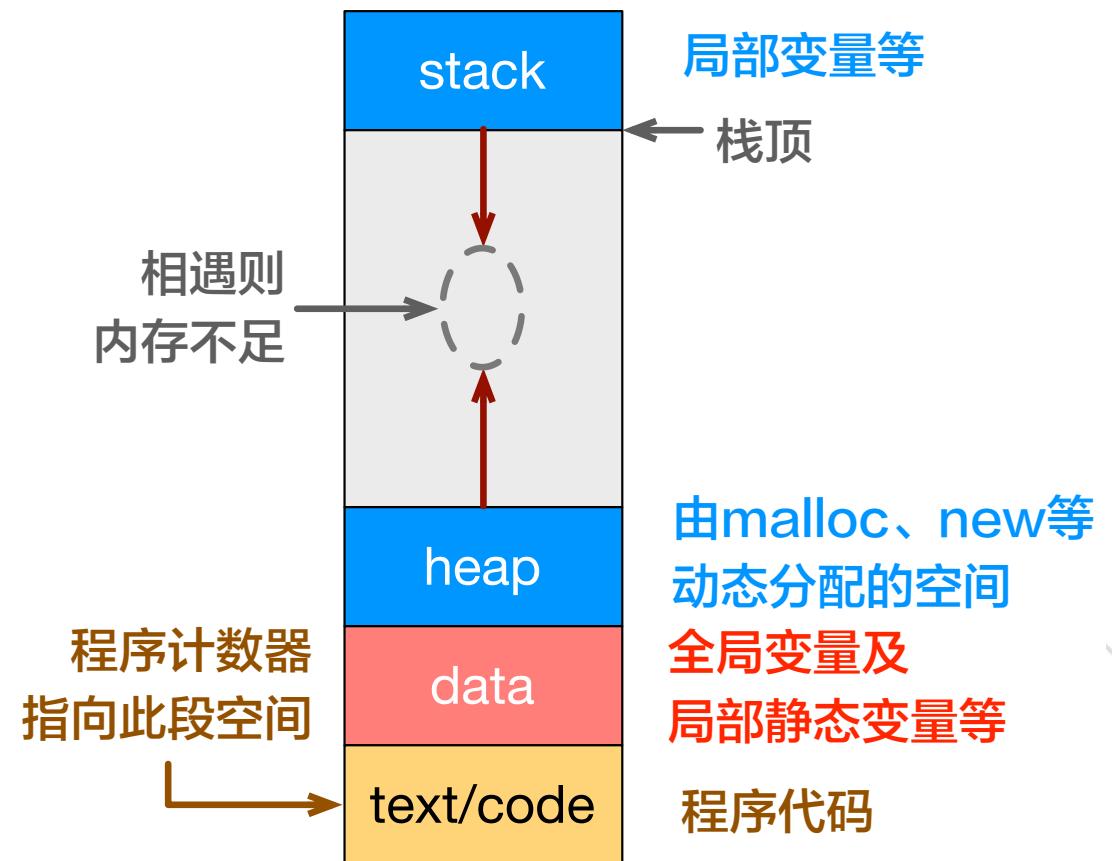
- 多个进程或线程“同时”进行
- 交错并发 (interleaved)：可通过系统调度由单核完成
- 并行 (parallel)：需多个运算核心同时完成
 - 多处理器 vs 单处理器多核
 - » 同一芯片上的多核通信速度更快
 - » 同一芯片上的多核能耗更低



● 多任务实现方式

– 进程：一个执行中的程序即为一个进程

- 每个进程有独立的程序计数器（program counter）、堆（heap）、栈（stack）、数据段（data section）、代码段（code section）等
- 程序运行状态由进程控制器（process control block）记录



PCB示例	
pointer	process state
process ID	
program counter	
registers	
scheduling info	
memory limits	
list of open files	
:	

进程状态：
new, running, waiting,
ready, terminated

寄存器：如累加器、堆栈指针等
用于进程中断后恢复

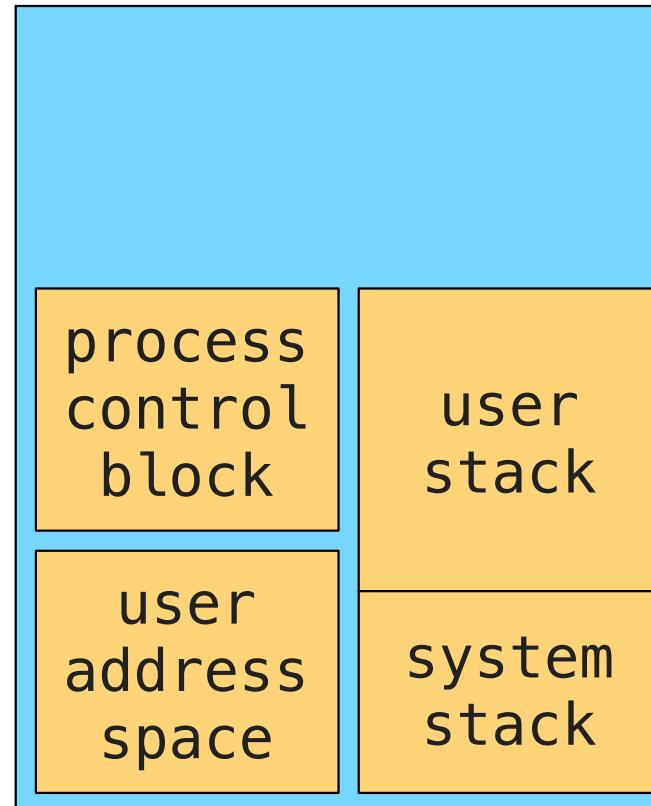
调度信息：如优先级等

● 多任务实现方式

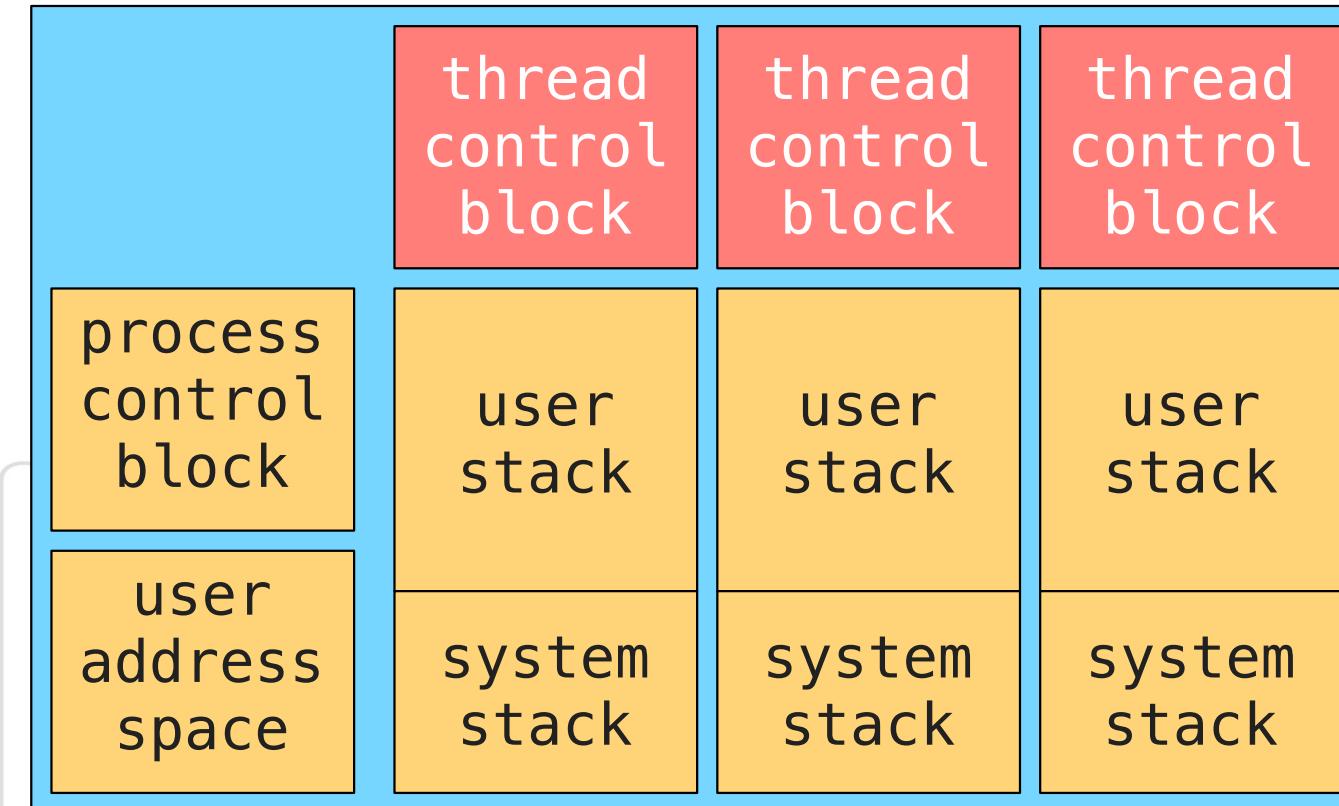
– 线程：常被称为轻量级进程（lightweight process）

- 与进程相似：每个线程有线程ID、程序计数器、寄存器、栈等
- 与进程不同：所有线程共享代码段、数据段及其他系统资源（如文件等）

单线程进程



多线程进程



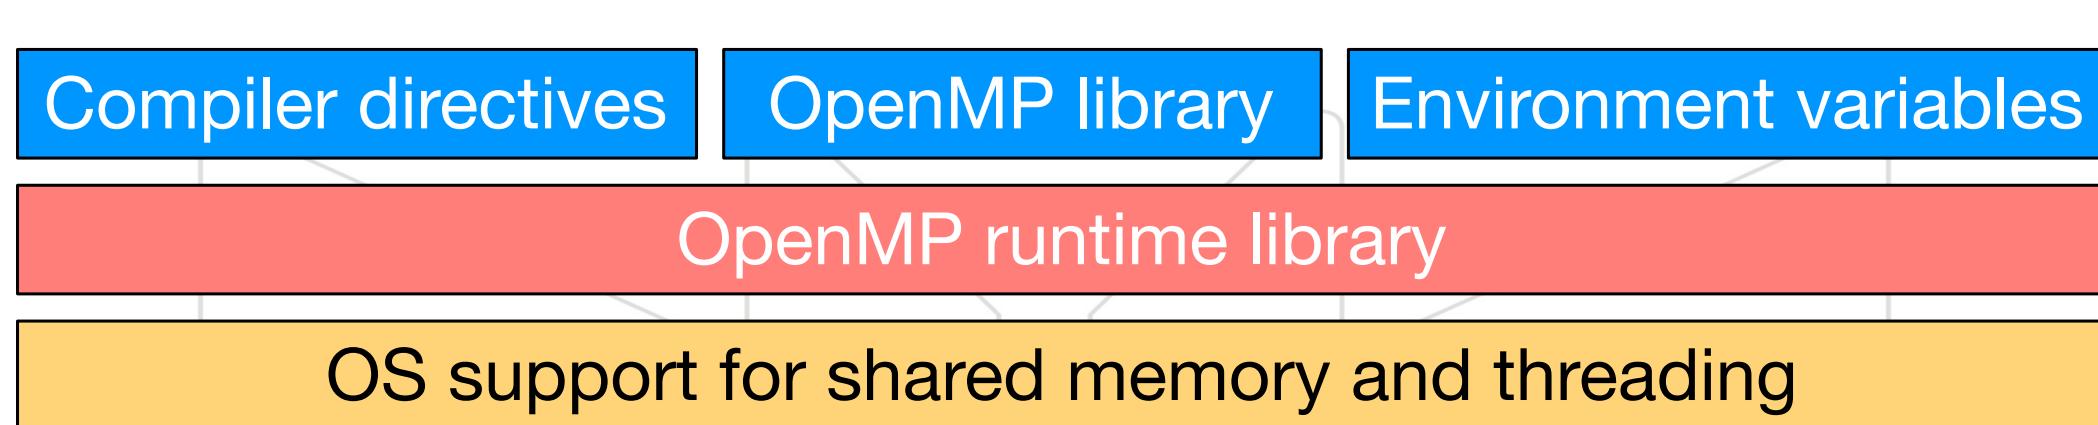
• OpenMP: Open Multi-Processing

– 多线程编程API

- 编译器指令 (#pragma)、库函数、环境变量
- 极大地简化了C/C++/Fortran多线程编程
- 并不是全自动并行编程语言
 - 其并行行为仍需由用户定义及控制

– 支持共享内存的多核系统

- 与CUDA、MPI所支持的硬件比较（讲义1）



- 什么是并行？
- 什么是OpenMP？
- 语法
- 同步机制
- 变量作用域
- 线程调度



● 预处理指令

- 设定编译器状态或指定编译器完成特定动作
 - 需要编译器支持相应功能
 - 否则将被忽略
- 举例: #pragma once
 - 指定头文件只被编译一次

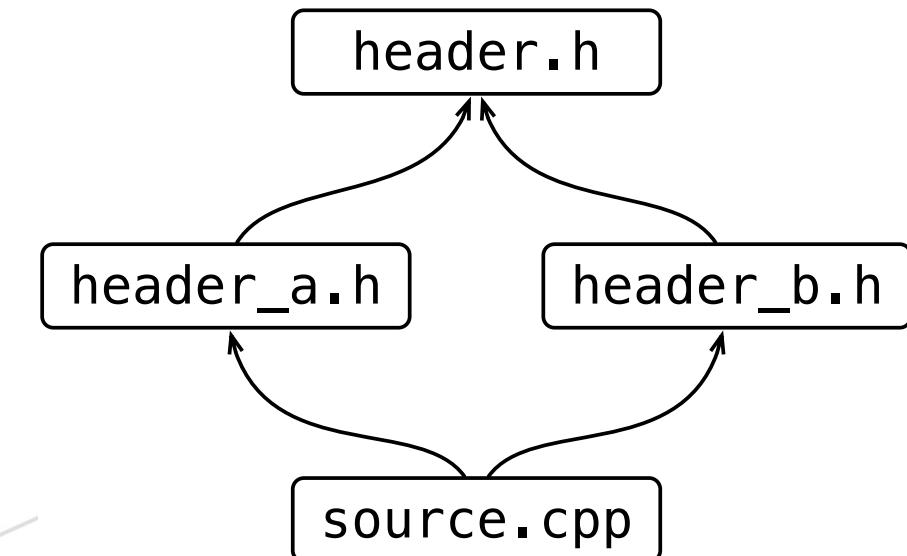
#pragma once

- 需要编译器支持
- 针对物理文件
- 需要用户保证头文件没有多份拷贝

#ifndef

- 不需要特定编译器
- 不针对物理文件
- 需要用户保证不同文件的宏名不重复

```
#ifndef HEADER_H
#define HEADER_H
...
#endif //HEADER_H
```



- 其他#pragma指令

- #pragma GCC poison printf
- #pragma warning (disable : 4996)

- OpenMP中的并行化声明由#pragma完成

- 格式为#pragma omp construct [clause [clause]...]
 - 如#pragma omp parallel for
 - 编译器如果不支持该指令则将直接忽略
- 其作用范围通常为一个代码区块

```
#pragma omp parallel for
for (int i=0; i<10; ++i){
    std::cout << i << std::endl;
}
```

● Windows

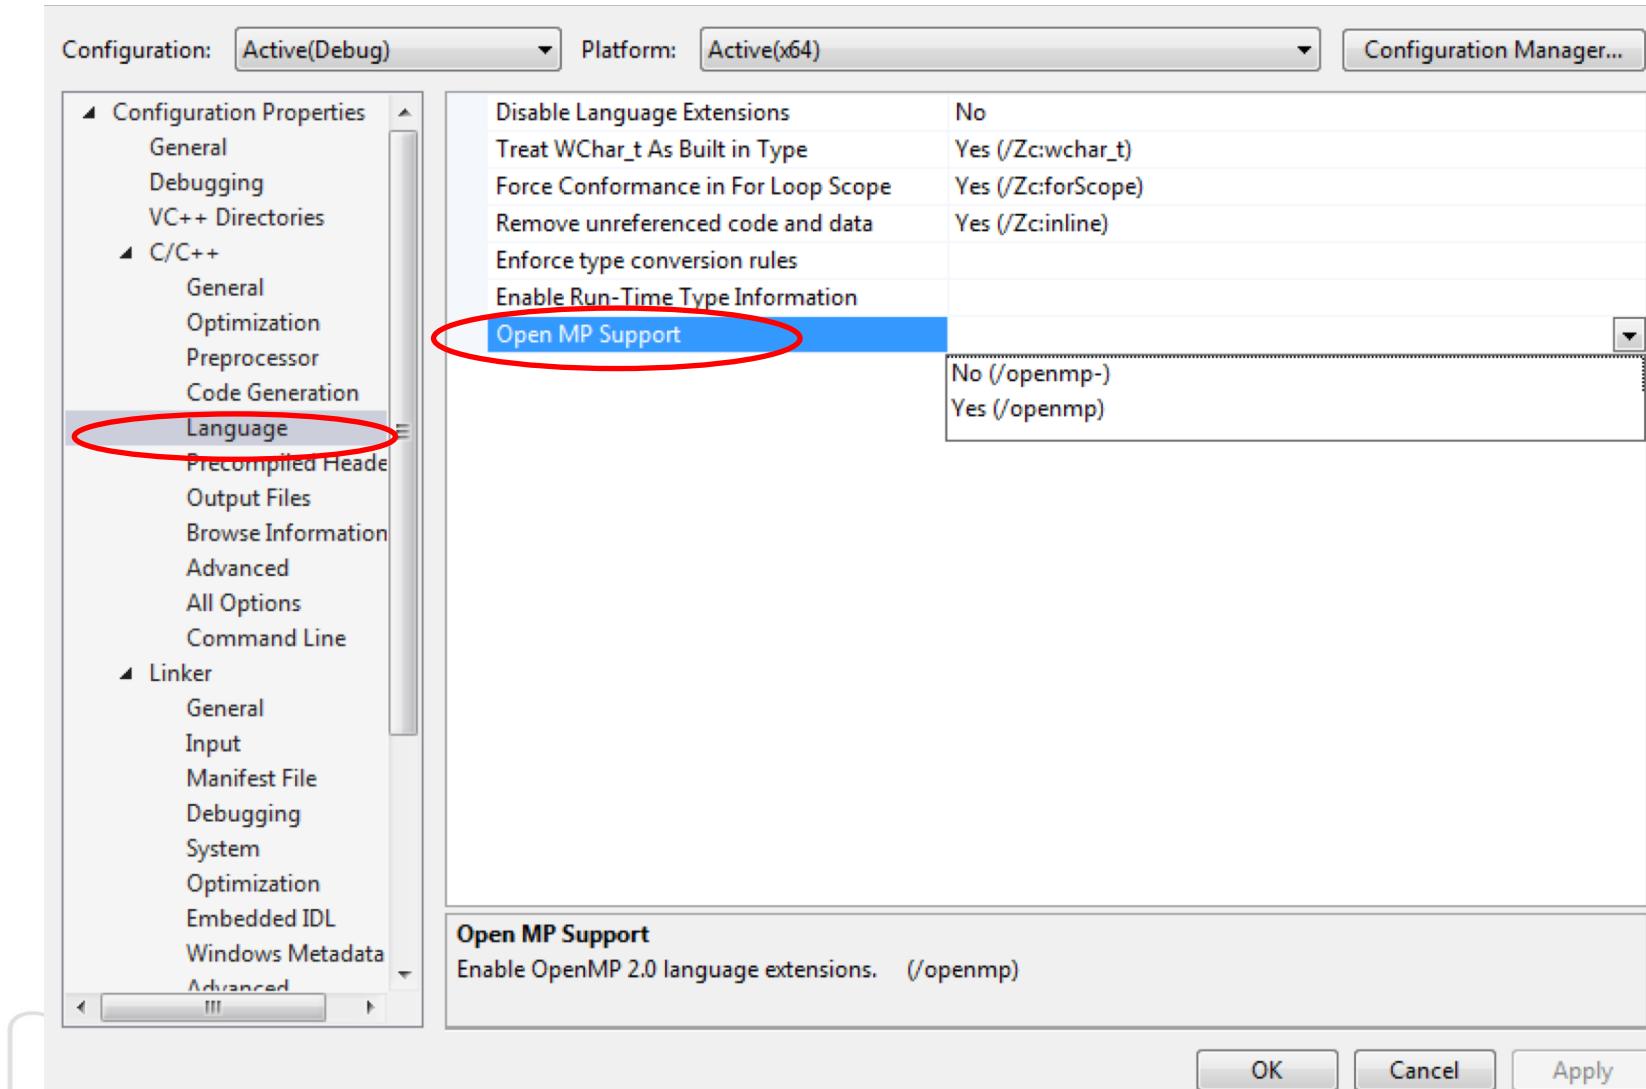
- 项目属性
 - >C/C++
 - >Language
 - >Open MP Support

● macOS/Linux

- 对于支持OpenMP的编译器
 - gcc: 在编译时增加
-fopenmp 标记

● 使用库函数

- `#include <omp.h>`



- 查看OpenMP版本

- 使用_OPENMP宏定义

```
#include <unordered_map>
#include <string>
#include <stdio>
#include <omp.h>

int main(int argc, char *argv[]) {
    std::unordered_map<unsigned, std::string> map{
        {200505, "2.5"}, {200805, "3.0"},
        {201107, "3.1"}, {201307, "4.0"},
        {201511, "4.5"}};
    printf("OpenMP version: %s.\n", map.at(_OPENMP).c_str());
    return 0;
}
```

编译: g++ -fopenmp openmp.cpp -o openmp_example

- 查看OpenMP版本
 - 使用__OPENMP宏定义

```
#include <unordered_map>
#include <string>
#include <cstdio>
#include <omp.h>

int main(int argc, char *argv[]) {
    std::unordered_map<unsigned, std::string> map{
        {200505, "2.5"}, {200805, "3.0"},
        {201107, "3.1"}, {201307, "4.0"},
        {201511, "4.5"}};
    printf("OpenMP version: %s.\n", map.at(__OPENMP).c_str());
    return 0;
}
```

成功编译运行（学院GPU集群）：
./openmp_example
OpenMP version: 4.5.

- 查看OpenMP版本
 - 使用__OPENMP宏定义

```
#include <unordered_map>
#include <string>
#include <cstdio>
#include <omp.h>

int main(int argc, char *argv[]) {
    std::unordered_map<unsigned, std::string> map{
        {200505, "2.5"}, {200805, "3.0"},
        {201107, "3.1"}, {201307, "4.0"},
        {201511, "4.5"}};
    printf("OpenMP version: %s.\n", map.at(__OPENMP).c_str());
    return 0;
}
```

- macOS默认编译器不支持OpenMP报错:
`clang: error: unsupported option '-fopenmp'`
- 解决方案 – 安装 llvm clang:
`brew install llvm
brew install libomp
echo 'export PATH="/usr/local/opt/llvm/bin:$PATH"' >> ~/.bash_profile`
- 编译:
`clang++ -fopenmp openmp.cpp -o openmp_example`

- 通过#pragma omp parallel 指明并行部分
- 无需改变串行代码

```
#include <stdio.h>
#include <omp.h>

int main()
{
#pragma omp parallel
{
    printf("Hello World\n");
}
return 0;
}
```

输出：
Hello World
Hello World

- 在输出中增加线程编号

- `omp_get_thread_num()`;

```
#include <stdio.h>
#include <omp.h>

int main()
{
#pragma omp parallel
{
    int thread = omp_get_thread_num();
    int max_threads = omp_get_max_threads();
    printf("Hello World (Thread %d of %d)\n", thread, max_threads);
}
return 0;
}
```

输出:

Hello World (Thread 0 of 8)
Hello World (Thread 4 of 8)
Hello World (Thread 1 of 8)
Hello World (Thread 7 of 8)
Hello World (Thread 3 of 8)
Hello World (Thread 2 of 8)
Hello World (Thread 6 of 8)
Hello World (Thread 5 of 8)

- 同一线程的多个语句是否连续执行？

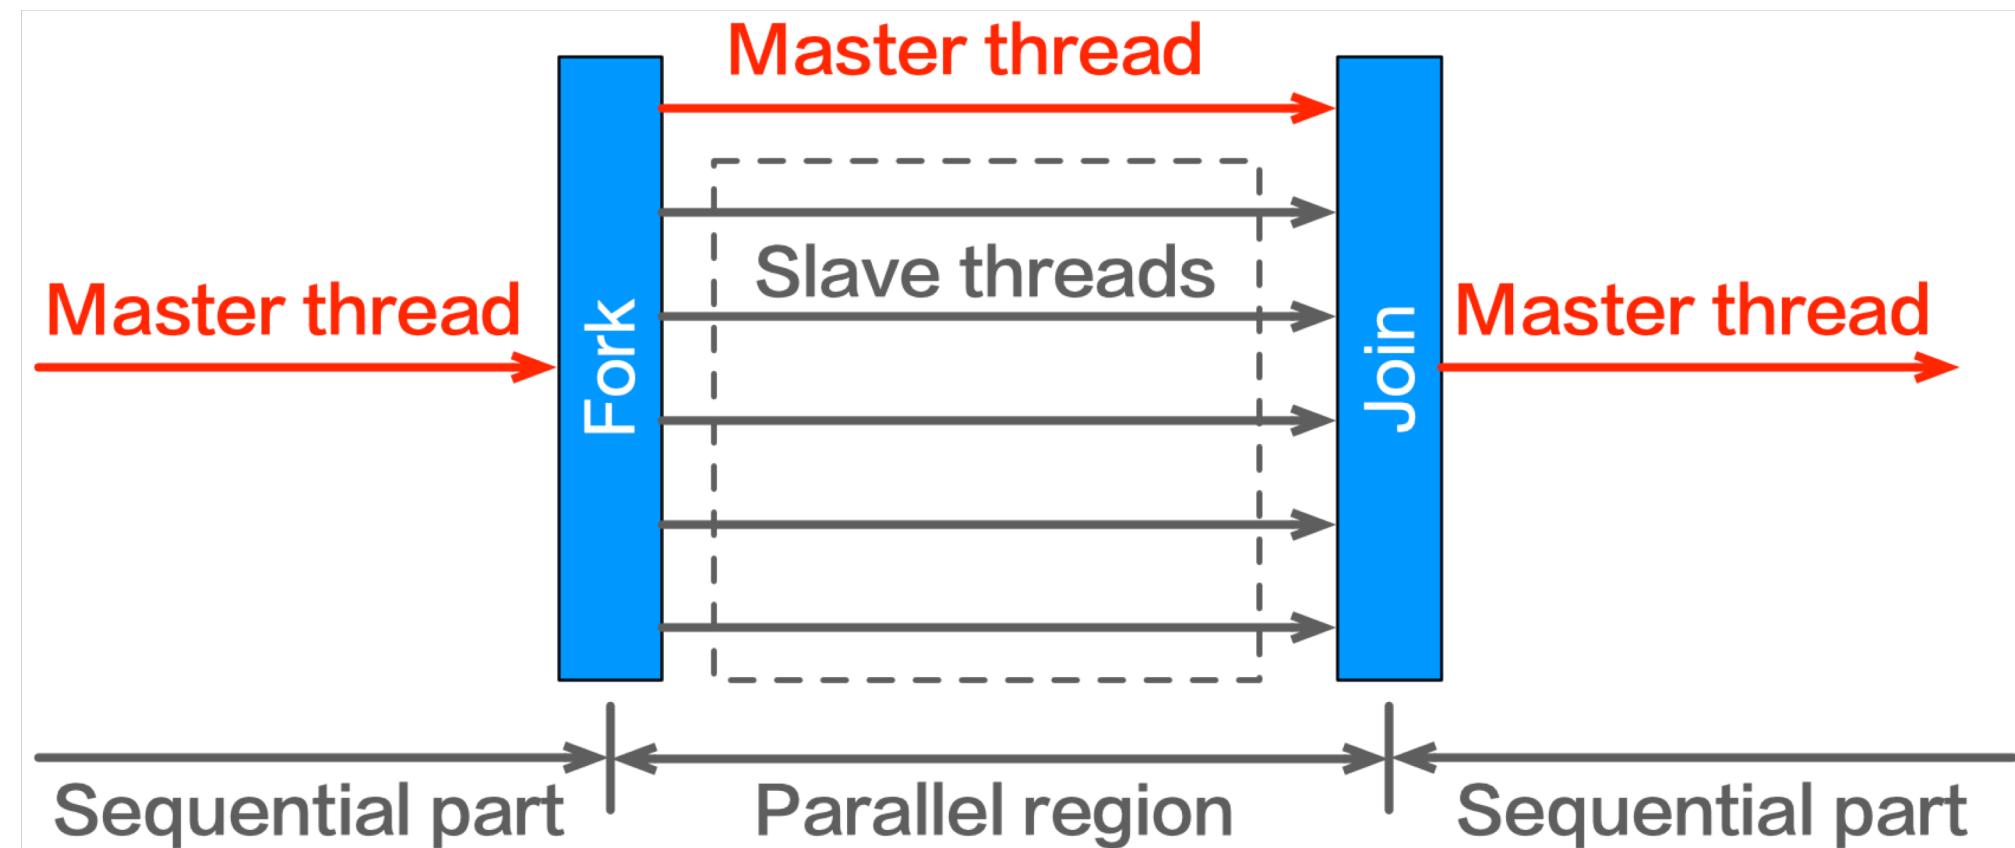
```
#include <stdio.h>
#include <omp.h>

int main()
{
#pragma omp parallel
{
    int thread = omp_get_thread_num();
    printf("hello(%d) ", thread);
    printf("world(%d) ", thread);
}
return 0;
}
```

输出：
hello(0) world(0)
hello(4) hello(1)
world(1) hello(7)
hello(3) world(7)
world(3) hello(6)
world(6) hello(5)
world(5) hello(2)
world(4) world(2)

- 使用分叉 (fork) 与交汇 (join) 模型

- Fork: 由主线程 (master thread) 创建一组从线程 (slave threads)
 - 主线程编号永远为0 (thread 0)
 - 不保证执行顺序
- Join: 同步终止所有线程并将控制权转移回至主线程



- 什么是并行？
- 什么是OpenMP？
- 语法
- 同步机制
- 变量作用域
- 线程调度



● 编译器指令

- `#pragma omp construct [clause [clause]...]{structured block}`
- 指明并行区域及并行方式
- clause子句
 - 指明详细的并行参数
 - 控制变量在线程间的作用域
 - 显式指明线程数目
 - 条件并行

```
#pragma omp parallel num_threads(16) ←  
{  
    int thread = omp_get_thread_num();  
    int max_threads = omp_get_max_threads();  
    printf("Hello World (Thread %d of %d)\n", thread, max_threads);  
}
```

- `num_threads(int)`

- 用于指明线程数目
- 当没有指明时，将默认使用`OMP_NUM_THREADS`环境变量
 - 环境变量的值为系统运算核心数目（或超线程数目）
 - 可以使用`omp_set_num_threads(int)`修改全局默认线程数
 - 可使用`omp_get_num_threads()`获取当前设定的默认线程数
 - `num_threads(int)`优先级高于环境变量
- `num_threads(int)`不保证创建指定数目的线程
 - 系统资源限制



● 并行for循环

- 将循环中的迭代分配到多个线程并行

```
#pragma omp parallel
{
    int n;
    for (n = 0; n < 4; n++){
        int thread = omp_get_thread_num();
        printf("thread %d \n", thread);
    }
}
```

输出是?

● 并行for循环

– 将循环中的迭代分配到多个线程并行

```
#pragma omp parallel
{
    int n;
    for (n = 0; n < 4; n++){
        int thread = omp_get_thread_num();
        printf("thread %d \n", thread);
    }
}
```

输出：

thread 3
thread 3
thread 0
thread 0
thread 0
thread 0
thread 0
thread 1
thread 1
thread 1
thread 1
thread 3
thread 3
thread 5
thread 5
...



● 并行for循环

– 将循环中的迭代分配到多个线程并行

- 风格1：在并行区域内加入#pragma omp for

```
#pragma omp parallel
{
    int n;
    #pragma omp for
    for (n = 0; n < 4; n++){
        int thread = omp_get_thread_num();
        printf("thread %d \n", thread);
    }
}
```

在并行区域内，for循环外
还可以加入其它并行代码

- 风格2：合并为#pragma omp parallel for

```
int n;
#pragma omp parallel for
for (n = 0; n < 4; n++) {
    int thread = omp_get_thread_num();
    printf("thread %d \n", thread);
}
```

写法更简洁

● 并行for循环

– 将循环中的迭代分配到多个线程并行

- 风格1：在并行区域内加入#pragma omp for

```
#pragma omp parallel
{
    int n;
    #pragma omp for
    for (n = 0; n < 4; n++){
        int thread = omp_get_thread_num();
        printf("thread %d \n", thread);
    }
}
```

输出：
thread 0
thread 2
thread 3
thread 1

- 风格2：合并为#pragma omp parallel for

```
int n;
#pragma omp parallel for
for (n = 0; n < 4; n++) {
    int thread = omp_get_thread_num();
    printf("thread %d \n", thread);
}
```

- OpenMP中的每个线程同样可以被并行化为一组线程

- OpenMP默认关闭嵌套

- 需要使用**omp_set_nested(1)**打开

```
omp_set_nested(1);

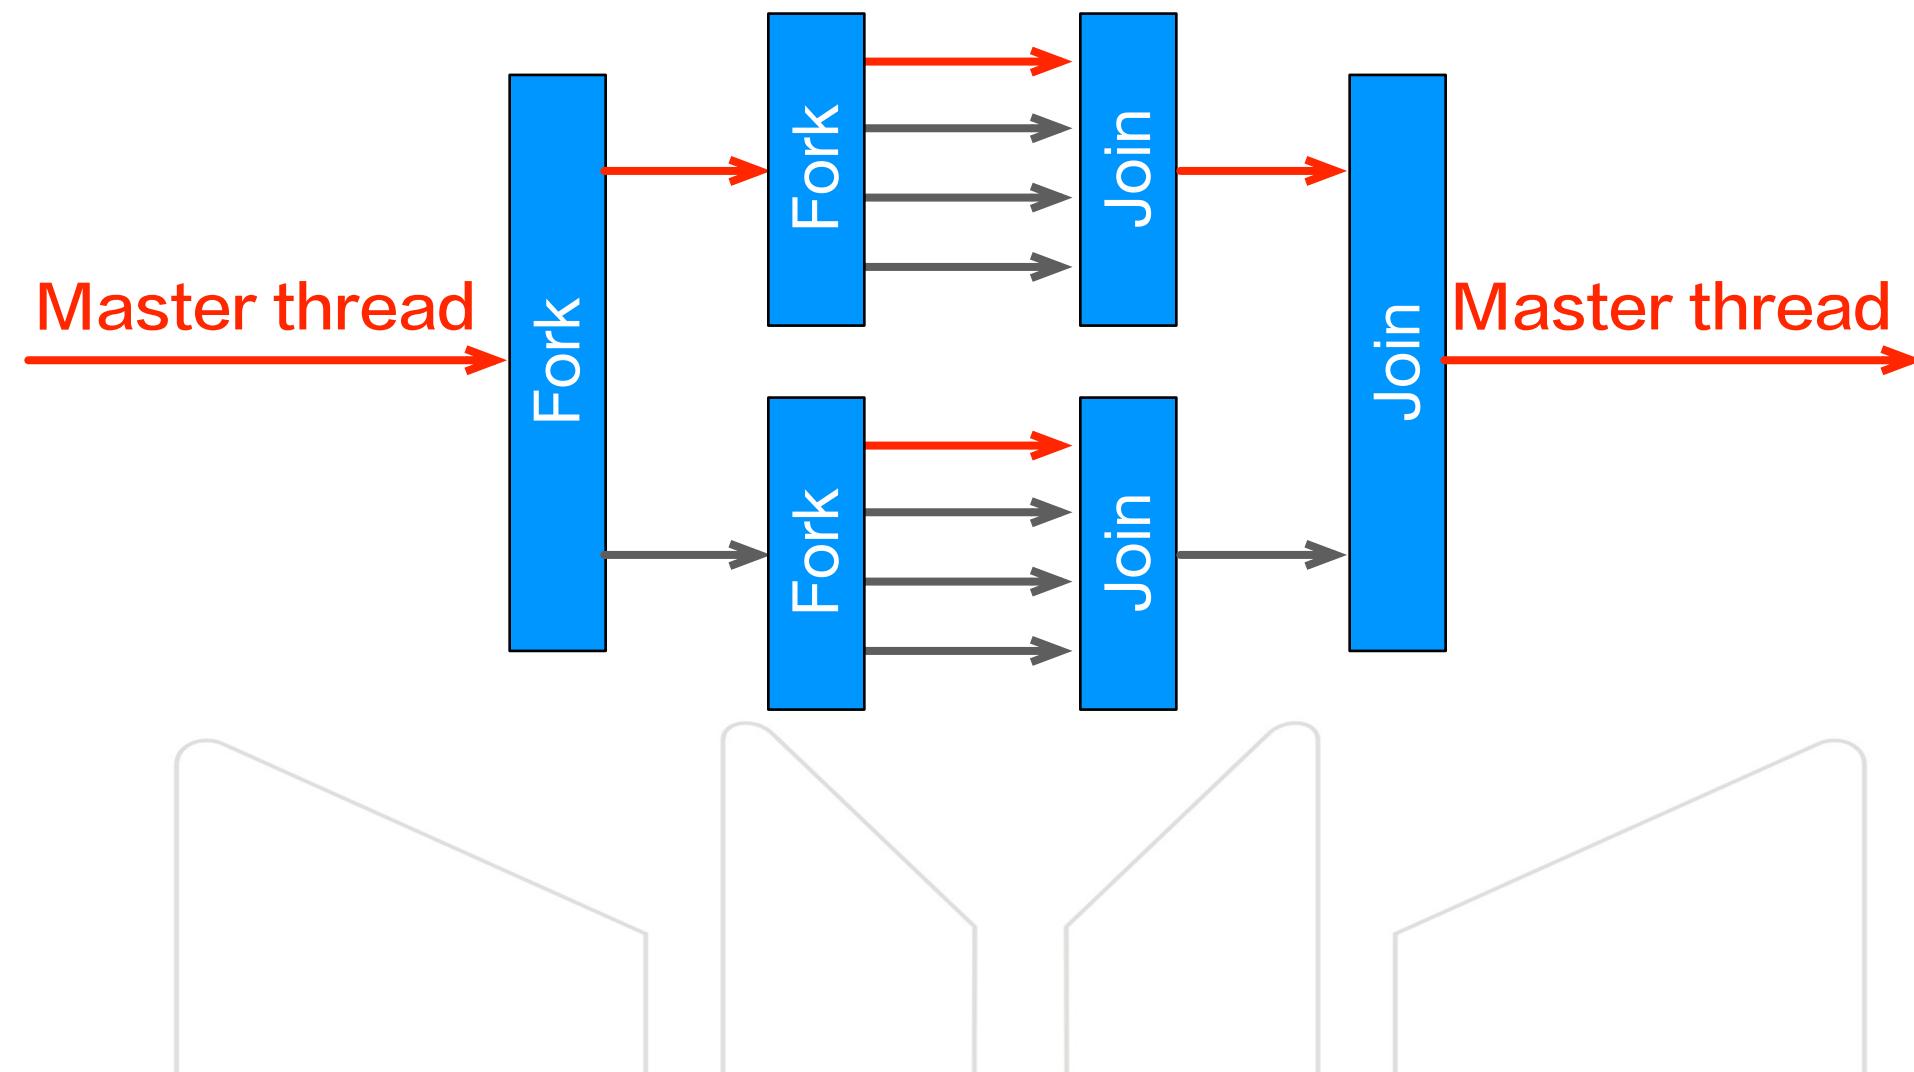
#pragma omp parallel for
for (int i = 0; i < 2; i++){
    int outer_thread = omp_get_thread_num();

#pragma omp parallel for
    for (int j = 0; j < 4; j++){
        int inner_thread = omp_get_thread_num();
        printf("Hello World (i = %d j = %d)\n",
               outer_thread, inner_thread);
    }
}
```

输出：

```
Hello World (i = 0 j = 0)
Hello World (i = 1 j = 0)
Hello World (i = 0 j = 2)
Hello World (i = 0 j = 3)
Hello World (i = 1 j = 1)
Hello World (i = 1 j = 2)
Hello World (i = 1 j = 3)
Hello World (i = 0 j = 1)
```

- OpenMP中的每个线程同样可以被并行化为一组线程
 - 仍然使用fork and join



● 语法限制

– 不能使用 != 作为判断条件

- `for (int i = 0; i!=8; ++i){`
- **error:** condition of OpenMP for loop must be a relational comparison ('<', '<=', '>', or '>=') of loop variable 'i'

– 循环必须为单入口单出口

- 不能使用 break、goto 等跳转语句
- **error:** 'break' statement cannot be used in OpenMP for loop

– (以上错误提示来自OpenMP 3.1)

● 数据依赖性

– 循环迭代相关 (loop-carried dependence)

- 依赖性与循环相关，去除循环则依赖性不存在

– 非循环迭代相关 (loop-independent dependence)

- 依赖性与循环无关，去除循环依赖性仍然存在

```
for (i = 1; i < n; i++) {  
    S1: a[i] = a[i - 1] + 1;  
    S2: b[i] = a[i];  
}
```

S1[i] → S1[i+1]: 循环相关
S1[i] → S2[i]: 循环无关

```
for (i = 1; i < n; i++)  
    for (j = 1; j < n; j++)  
        S3: a[i][j] = a[i][j - 1] + 1;
```

S3[i, j] → S3[i, j+1]:
i循环无关, j循环相关

```
for (i = 1; i < n; i++)  
    for (j = 1; j < n; j++)  
        S4: a[i][j] = a[i - 1][j] + 1;
```

S4[i, j] → S4[i+1, j]:
i循环相关, j循环无关

- 什么是并行？
- 什么是OpenMP？
- 语法
- 同步机制
- 变量作用域
- 线程调度



- OpenMP是多线程共享地址架构
 - 线程可通过共享变量通信
- 线程及其语句执行具有不确定性
 - 共享数据可能造成竞争条件 (race condition)
 - 竞争条件：程序运行的结果依赖于不可控的执行顺序
- 必须使用同步机制避免竞争条件的出现
 - 同步机制将带来巨大开销
 - 尽可能改变数据的访问方式减少必须的同步次数



- 语句执行顺序造成结果不一致

```
int a[3] = { 3, 4, 5};
```

thread 1

```
a[1] = a[0] + a[1];
```

thread 2

```
a[2] = a[1] + a[2];
```

a = { 3, ?, ? }



- 语句执行顺序造成结果不一致

```
int a[3] = { 3, 4, 5};
```

thread 1

```
a[1] = a[0] + a[1];
```

thread 2

```
a[2] = a[1] + a[2];
```

$a = \{ 3, ?, ? \}$

- 先执行 thread 1 再执行 thread 2

- $a[1]=a[0]+a[1]=3+4=7$; $a[2]=a[1]+a[2]=7+5=12$;
 - $a = \{ 3, 7, 12 \}$

- 先执行 thread 2 再执行 thread 1

- $a[2]=a[1]+a[2]=4+5=9$; $a[1]=a[0]+a[1]=3+4=7$;
 - $a = \{ 3, 7, 9 \}$

- 高级语言的语句并非原子操作

```
int count=10;  
  
thread 1           thread 2  
count++;            count--;
```

count = 9, 10, 11?



- 高级语言的语句并非原子操作

```
int count=10;  
  
thread 1           thread 2  
LOAD Reg, count    LOAD Reg, count  
ADD #1              SUB #1  
STORE Reg, count   STORE Reg, count
```

count = 9, 10, 11?



- 高级语言的语句并非原子操作

int count=10;			
thread 1		thread 2	
	Reg	count	Reg
LOAD	10	10	
ADD	11	10	
	11	10	10 LOAD
	11	10	9 SUB
	11	9	9 STORE
STORE	11	11	

count = 11

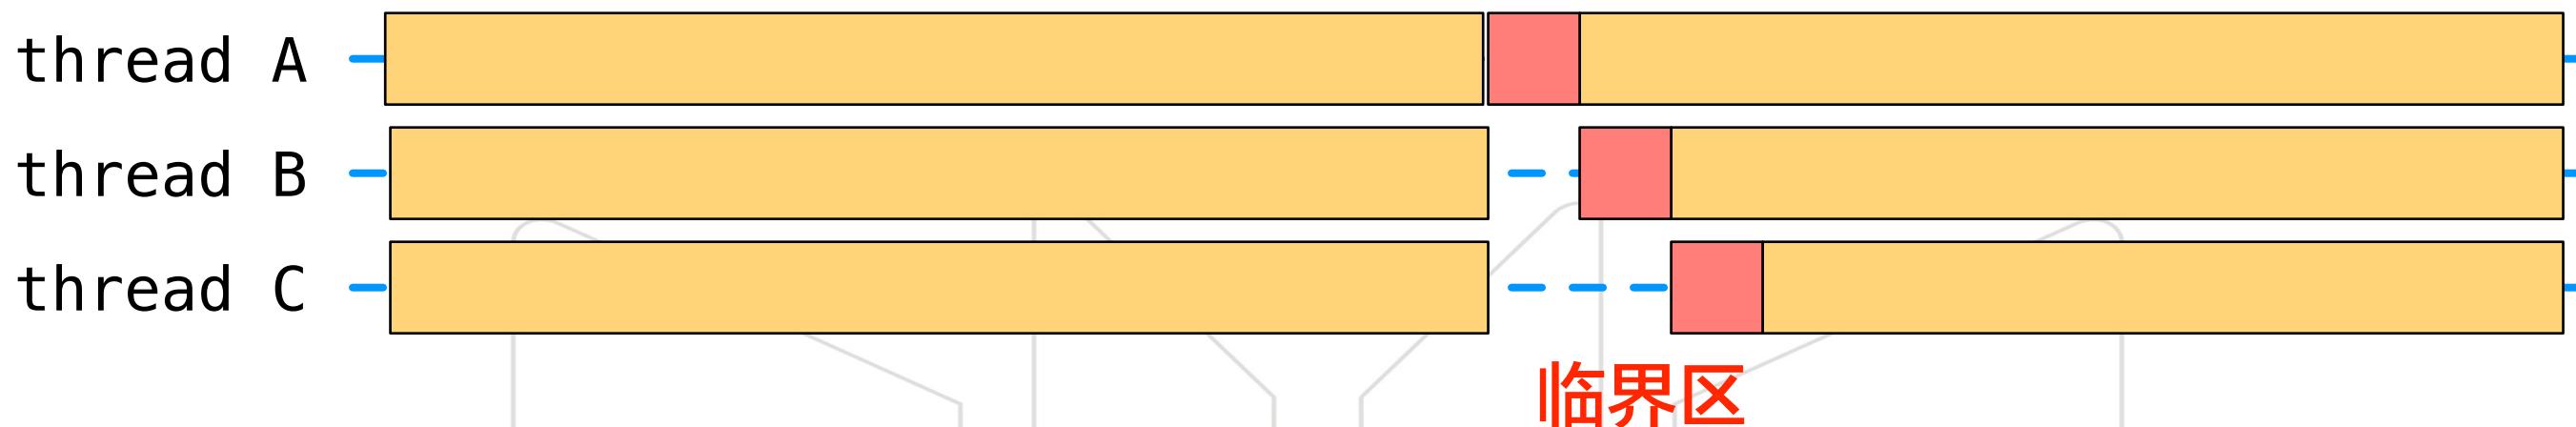
- 高级语言的语句并非原子操作

		int count=10;	
		thread 1	thread 2
	Reg	count	Reg
	LOAD	10	10
LOAD	10	10	10
ADD	11	10	10
	11	10	9
STORE	11	11	9
	11	9	9
			SUB
			STORE

count = 9

● 临界区 (critical section)

- `#pragma omp critical`
- 指的是一个访问共用资源（例如：共用设备或是共用存储器）的程序片段，而这些共用资源又无法同时被多个线程访问的特性
 - 同一时间内只有一个线程能执行临界区内代码
 - 其他线程必须等待临界区内线程执行完毕后才能进入临界区
 - 常用来保证对共享数据的访问之间互斥



● 临界区 (critical section)

- `#pragma omp critical`
- 比照操作系统中信号量 (semaphore) 与P、V操作

```
#pragma omp critical
{
    ...
    critical section;
    ...
}
```

```
Semaphore a;
P(a);
...
critical section;
...
V(a);
```



● 临界区 (critical section)

– 举例：统计随机数分布

- 随机产生1000个[0-20)之间的整数
- 统计每个数字出现频率

无临界区：

```
#pragma omp parallel for
for(int i=0; i<1000; ++i){
    int value = rand()%20;
    histogram[value]++;
}

int total = 0;
for(int i = 0; i < 20; i++){
    total += histogram[i];
    cout<<histogram[i]<<" ";
}
cout<<endl<<"total: "<<total<<endl;
```

输出：

```
25 31 26 34 40 47 24 29 44 44
31 26 41 38 32 45 26 54 45 27
total: 709
```

● 临界区 (critical section)

– 举例：统计随机数分布

- 随机产生1000个[0-20)之间的整数
- 统计每个数字出现频率

有临界区：

```
#pragma omp parallel for
for(int i=0; i<1000; ++i){
    int value = rand()%20;

    #pragma omp critical
    {
        histogram[value]++;
    }
}
```

输出：

```
60 47 28 54 52 50 33 56 44 53
61 58 43 47 52 54 50 52 53 53
total: 1000
```

- 原子 (atomic) 操作

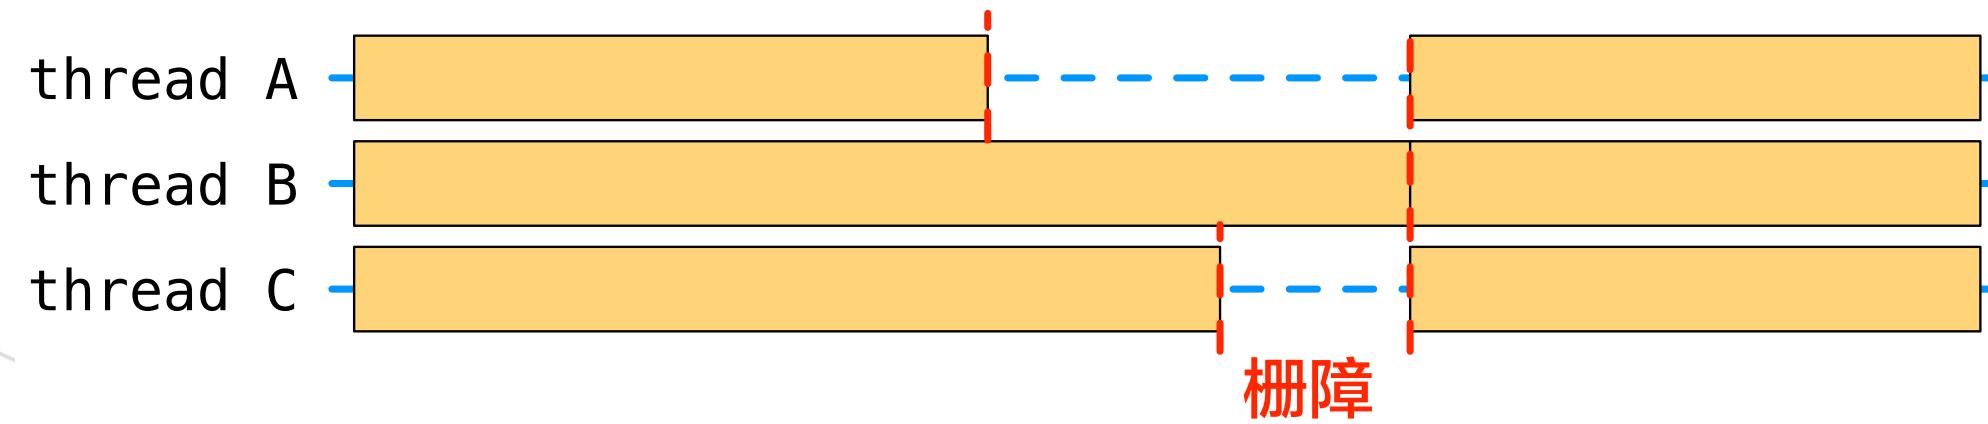
- `#pragma omp atomic`
- 保证下一语句在同一时间只能被一个线程执行
 - 常用来做计数器、求和等
- 原子操作通常比临界区执行更快
- 临界区的作用范围更广，能够实现的功能更复杂

```
#pragma omp parallel for
for(int i=0; i<1000; ++i){
    int value = rand()%20;
    #pragma omp atomic
    histogram[value]++;
}
```

● 栅障 (barrier)

- `#pragma omp barrier`
- 在栅障点处同步所有线程
 - 先运行至栅障点处的线程必须等待其他线程
 - 常用来等待某部分任务完成再开始下一部分任务
 - 每个并行区域的结束点默认自动同步线程

```
#pragma omp parallel
{
    function_A()
    #pragma omp barrier
    function_B();
}
```



● 栅障 (barrier)

- 并行随机数统计及并行求和

```
int total = 0;  
  
#pragma omp parallel num_threads(20)  
{  
    for(int i=0; i<50; ++i){  
        int value = rand()%20;  
        #pragma omp atomic  
        histogram[value]++;  
    }  
  
    int thread = omp_get_thread_num();  
    #pragma omp atomic  
    total += histogram[thread];  
}
```

输出:
total: 619

← 求和时可能其他线程还没完成统计

● 栅障 (barrier)

- 并行随机数统计及并行求和

```
int total = 0;  
  
#pragma omp parallel num_threads(20)  
{  
    for(int i=0; i<50; ++i){  
        int value = rand()%20;  
        #pragma omp atomic  
        histogram[value]++;  
    }  
    #pragma omp barrier ← 使用栅障同步线程  
    int thread = omp_get_thread_num();  
    #pragma omp atomic  
    total += histogram[thread];  
}
```

输出：
total: 1000

● 栅障 (barrier)

- 并行随机数统计及并行求和
 - 这两段代码结果是否相同？

```
int total = 0;

#pragma omp parallel num_threads(20)
{
    for(int i=0; i<50; ++i){
        int value = rand()%20;
        #pragma omp atomic
        histogram[value]++;
    }
    #pragma omp barrier
    int thread = omp_get_thread_num();
    #pragma omp atomic
    total += histogram[thread];
}
```

```
int total = 0;

#pragma omp parallel num_threads(20)
{
    #pragma omp for
    for(int i=0; i<1000; ++i){
        int value = rand()%20;
        #pragma omp atomic
        histogram[value]++;
    }
    int thread = omp_get_thread_num();
    #pragma omp atomic
    total += histogram[thread];
}
```

- **#pragma omp single {}**

- 用于保证{}内的代码由一个线程完成
- 常用于输入输出或初始化
- 由第一个执行至此处的线程执行
- 同样会产生一个隐式栅障
 - 可由**#pragma omp single nowait**去除

- **#pragma omp master {}**

- 与single相似，但指明由主线程执行
- 与使用IF的条件并行等价
 - **#pragma omp parallel IF(omp_get_thread_num() == 0) nowait**
 - 默认不产生隐式栅障

● **#pragma omp master {}**

– 在下面代码中与atomic等价

```
int total = 0;

#pragma omp parallel
{
    #pragma omp for
    for(int i=0; i<1000; ++i){
        int value = rand()%20;
        #pragma omp atomic
        histogram[value]++;
    }
    #pragma omp master
    {
        for(int i=0; i<20; ++i){
            total += histogram[i];
        }
    }
}
```

```
int total = 0;

#pragma omp parallel num_threads(20)
{
    #pragma omp for
    for(int i=0; i<1000; ++i){
        int value = rand()%20;
        #pragma omp atomic
        histogram[value]++;
    }

    int thread = omp_get_thread_num();
    #pragma omp atomic
    total += histogram[thread];
}
```

- 指明如何将线程局部结果汇总

- 如#**pragma omp for reduction(+: total)**
- 支持的操作: +, -, *, &, |, && and ||

```
int total = 0;

#pragma omp parallel num_threads(20)
{
    for(int i=0; i<50; ++i){
        int value = rand()%20;
        #pragma omp atomic
        histogram[value]++;
    }
    #pragma omp barrier
    int thread = omp_get_thread_num();
    #pragma omp atomic
    total += histogram[thread];
}
```

```
int total = 0;

#pragma omp parallel
{
    #pragma omp for
    for(int i=0; i<1000; ++i){
        int value = rand()%20;
        #pragma omp atomic
        histogram[value]++;
    }
    #pragma omp for reduction(+: total) ←
    for(int i=0; i<20; ++i){
        total += histogram[i];
    }
}
```

- 什么是并行？
- 什么是OpenMP？
- 语法
- 同步机制
- 变量作用域
- 线程调度



- OpenMP与串行程序的作用域不同
 - OpenMP中必须指明变量为**shared**或**private**
 - Shared: 变量为所有线程所共享
 - 并行区域外定义的变量默认为**shared**
 - Private: 变量为线程私有, 其他线程无法访问
 - 并行区域内定义的变量默认为**private**
 - 循环计数器默认为**private**



- Shared 与 private

```
int histogram[20]; ← shared
init_histogram(histogram);

int total = 0; ← shared

#pragma omp parallel
{
    int i; ← 循环计数器为private!
    #pragma omp for
    for(i=0; i<1000; ++i){
        private → int value = rand()%20;
        #pragma omp atomic
        histogram[value]++;
    }
}
```

● 显式作用域定义

- 显式指明变量的作用域
- shared (var)
 - 指明变量var为shared
- default (none/shared/private)
 - 指明变量的默认作用域
 - 如果为none则必须指明并行区域内每一变量的作用域

```
int a, b = 0, c;  
#pragma omp parallel default(None) shared(b)  
{  
    b += a;  
}
```

error: variable 'a' must have explicitly specified data sharing attributes

● 显式作用域定义

- **private (var)**
 - 指明变量var为private

```
int i = 10;
#pragma omp parallel for private(i)
for (int j=0; j<4; ++j) {
    printf("Thread %d: i = %d\n", omp_get_thread_num(), i);
}
printf("i = %d\n", i);
```

输出:

```
Thread 0: i = 1
Thread 1: i = 0
Thread 3: i = 0
Thread 2: i = 0
i = 10
```

- **firstprivate(var)**

- 指明变量var为private , 同时表明该变量使用master thread中变量值初始化

```
int i = 10;
#pragma omp parallel for firstprivate(i)
for (int j=0; j<4; ++j) {
    printf("Thread %d: i = %d\n", omp_get_thread_num(), i);
}
printf("i = %d\n", i);
```

输出:

```
Thread 0: i = 10
Thread 3: i = 10
Thread 2: i = 10
Thread 1: i = 10
i = 10
```

● 显式作用域定义

- **private (var)**
 - 指明变量var为private

```
int i = 10;
#pragma omp parallel for private(i)
for (int j=0; j<4; ++j) {
    printf("Thread %d: i = %d\n", omp_get_thread_num(), i);
}
printf("i = %d\n", i);
```

输出:

```
Thread 0: i = 1
Thread 1: i = 0
Thread 3: i = 0
Thread 2: i = 0
i = 10
```

- **lastprivate(var)**

- 指明变量var为private , 同时表明结束后一层迭代将结果赋予该变量

```
int i = 10;
#pragma omp parallel for lastprivate(i)
for (int j=0; j<4; ++j) {
    printf("Thread %d: i = %d\n", omp_get_thread_num(), i);
}
printf("i = %d\n", i);
```

输出:

```
Thread 0: i = 1
Thread 3: i = 0
Thread 1: i = 0
Thread 2: i = 0
i = 0
```

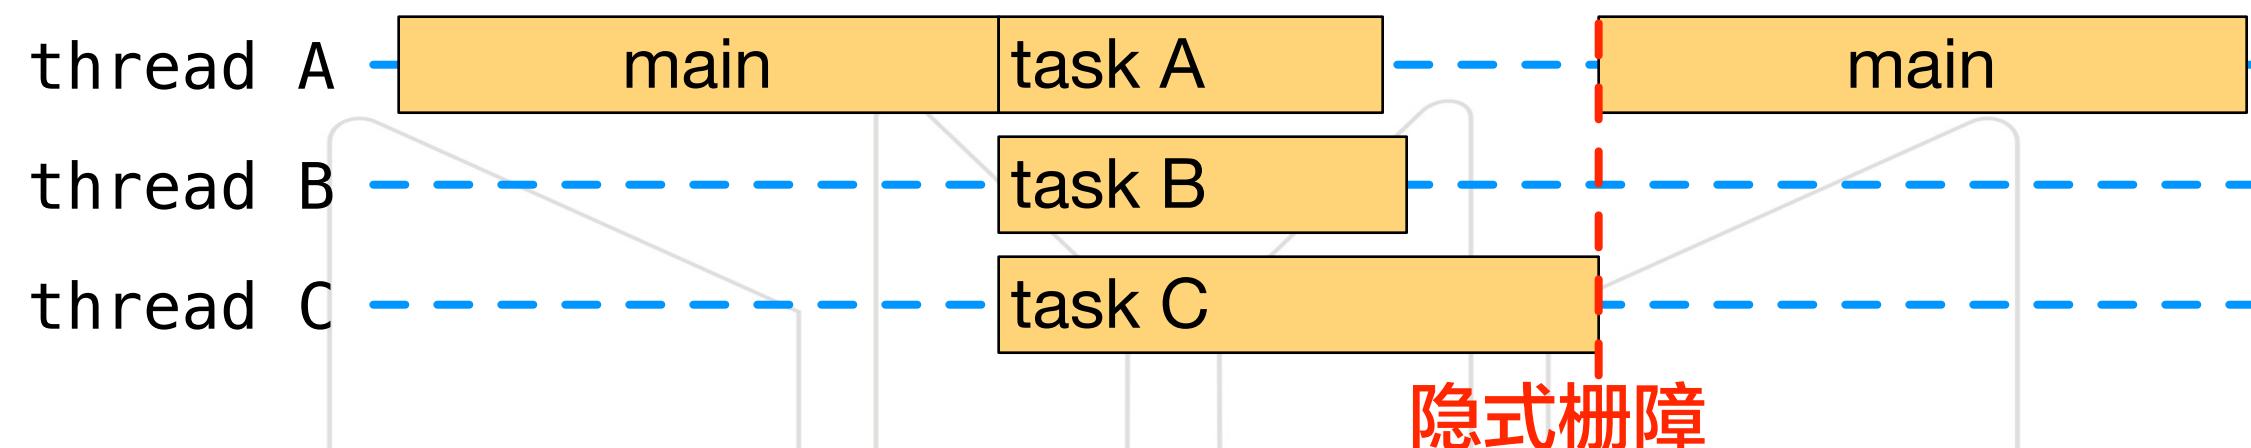
● 数据并行

- 同样指令作用在不同数据上
- 前述例子均为数据并行

● 任务并行

- 线程可能执行不同任务
- `#pragma omp sections`
- 每个section由一个线程完成
- 同样有隐式栅障（可使用nowait去除）

```
#pragma omp parallel
#pragma omp sections
{
    #pragma omp section
task_A();
    #pragma omp section
task_B();
    #pragma omp section
task_C();
}
```



- 什么是并行？
- 什么是OpenMP？
- 语法
- 同步机制
- 变量作用域
- 线程调度



- 当迭代数多于线程数时，需要调度线程
 - 某些线程将执行多个迭代
 - `#pragma omp parallel for schedule(type,[chunk size])`
 - type 包括 static, dynamic, guided, auto, runtime
 - 默认为static

```
#pragma omp parallel for num_threads(4)
for (int i=0; i<6; ++i)
{
    int thread = omp_get_thread_num();
    printf("thread %d\n", thread);
}
```

输出：
thread 1
thread 1
thread 3
thread 0
thread 0
thread 2

● Static调度

- 调度由编译器静态决定
- `#pragma omp parallel for schedule(type, [chunk size])`
 - 每个线程轮流获取 chunk size 个迭代任务
 - 默认chunk size 为 $n/threads$

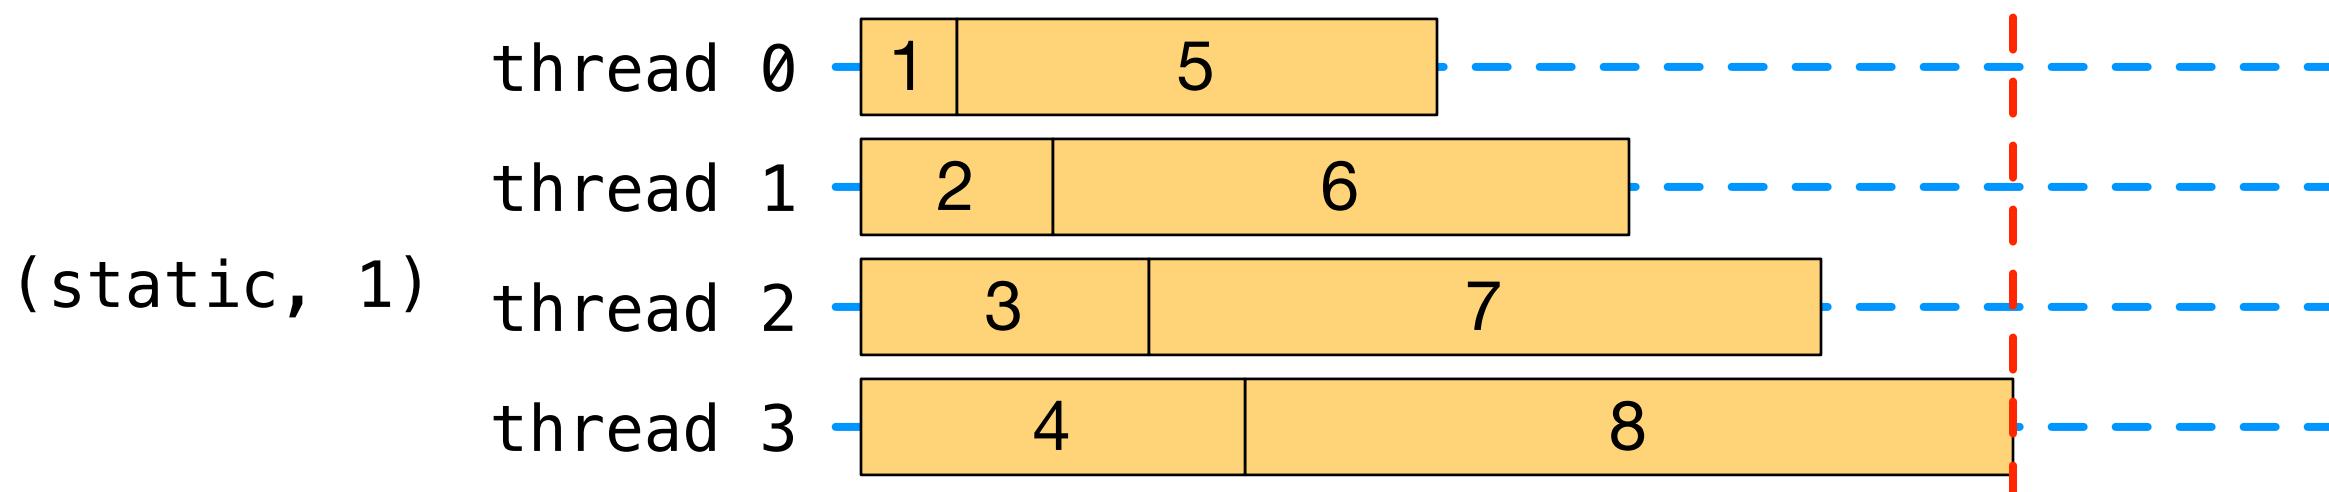
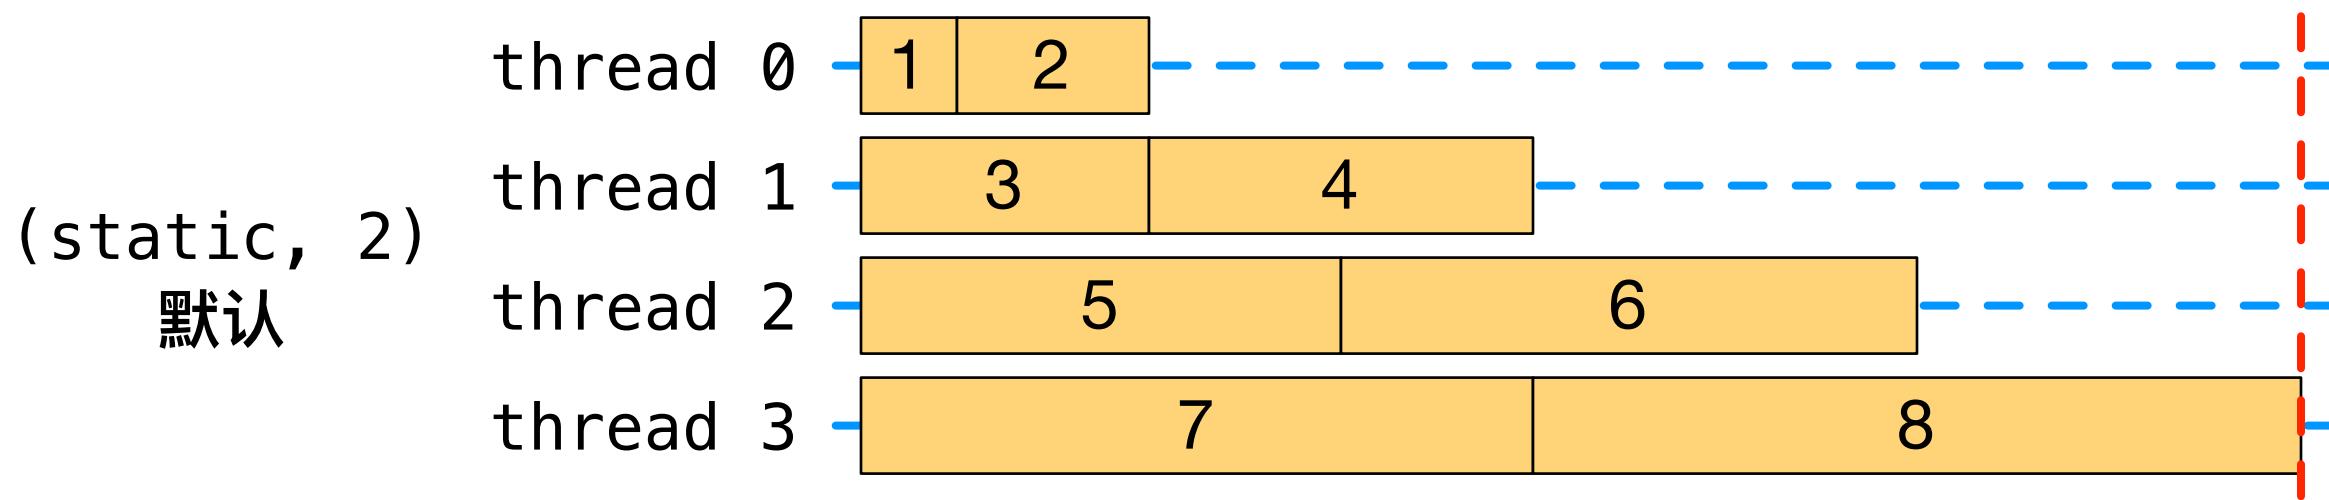
	(static, 1)			
thread 0	0	4	8	12
thread 1	1	5	9	13
thread 2	2	6	10	14
thread 3	3	7	11	15

	(static, 2)			
thread 0	0	1	8	9
thread 1	2	3	10	11
thread 2	4	5	12	13
thread 3	6	7	14	15

	(static, 4)			
thread 0	0	1	2	3
thread 1	4	5	6	7
thread 2	8	9	10	11
thread 3	12	13	14	15

● Static调度

– 线程的负载可能不均匀



● Dynamic调度

- 在运行中动态分配任务
- 迭代任务依然根据chunk size划分成块
- 线程完成一个chunk后向系统请求下一个chunk

● Guided调度

- 与dynamic类似
- 但分配的chunk大小在运行中递减
 - 最小不能小于chunk size参数

● Auto 与 runtime

- “Note that keywords auto and runtime aren't adequate.”

- 当#pragma指令无法为编译器理解时
 - 不会报错！
 - 错在哪儿？
 - #pragma omp parallel
- 参考OpenMP的32个常见陷阱
 - <https://software.intel.com/en-us/articles/32-openmp-traps-for-c-developers>



● OpenMP Reference Guide

- <https://www.openmp.org/wp-content/uploads/OpenMP-4.5-1115-CPP-web.pdf>

OpenMP API 4.5 C/C++

Page 1

OpenMP® C/C++

OpenMP 4.5 API C/C++ Syntax Reference Guide

OpenMP Application Program Interface (API) is a portable, scalable model that gives parallel programmers a simple and flexible interface for developing portable parallel applications. OpenMP supports multi-platform shared-memory parallel programming in C/C++ and Fortran on all architectures, including Unix platforms and Windows platforms. See www.openmp.org for specifications.

- Text in this color indicates functionality that is new or changed in the OpenMP API 4.5 specification.
- [n.n.n] Refers to sections in the OpenMP API 4.5 specification.
- [n.n.n] Refers to sections in the OpenMP API 4.0 specification.

Directives and Constructs for C/C++

An OpenMP executable directive applies to the succeeding structured block or an OpenMP construct. Each directive starts with `#pragma omp`. The remainder of the directive follows the conventions of the C and C++ standards for compiler directives. A *structured-block* is a single statement or a compound statement with a single entry at the top and a single exit at the bottom.

parallel [2.5] [2.5]
Forms a team of threads and starts parallel execution.
#pragma omp parallel [clause[,]clause] ...
 structured-block
clause:
 `if[/ parallel :] scalar-expression`
 `num_threads(integer-expression)`
 `default(shared | none)`
 `private(list)`
 `firstprivate(list)`
 `shared(list)`
 `copyin(list)`
 `reduction(reduction-identifier: list)`
 `proc_bind(master | close | spread)`

sections [2.7.2] [2.7.2]
A noniterative worksharing construct that contains a set of structured blocks that are to be distributed among and executed by the threads in a team.
#pragma omp sections [clause[,] clause] ...
 {
 `[#pragma omp section]`
 structured-block
 `[#pragma omp section]`
 structured-block
 ...
 }
clause:
 `private(list)`
 `firstprivate(list)`

for simd [2.8.3] [2.8.3]
Specifies that a loop that can be executed concurrently using SIMD instructions, and that those iterations will also be executed in parallel by threads in the team.
#pragma omp for simd [clause[,]clause] ...
 for-loops
clause:
 Any accepted by the `simd` or `for` directives with identical meanings and restrictions.

task [2.9.1] [2.11.1]
Defines an explicit task. The data environment of the task is created according to data-sharing attribute clauses on task construct and any defaults that apply.

● 软硬件环境

- CPU多线程并行库
 - 编译器指令、库函数、环境变量
- 共享内存的多核系统

● 基本语法

- `#pragma omp construct [clause [clause]...]{structured block}`
- 指明并行区域: `#pragma omp parallel`
- 循环: `#pragma omp (parallel) for`
- 嵌套: `omp_set_nested(1)`
- 常用函数: `omp_get_thread_num(); num_threads(int);`
- 同步: `#pragma omp critical/atomic/barrier、nowait`
- 变量作用域: `default(none/shared/private), shared(), private(), firstprivate(), last private()`
- 调度: `schedule(static/dynamic/guided, [chunk_size])`

Questions?

